

Opponent-Aware Q-Learning for Leduc Hold'em Poker

Sparsh Agrawal
BT2024084

Tejas Kollipara
BT2024147

Varun E
BT2024220

Divy Dobariya
BT2024225

Abstract—We investigate opponent-aware Q-learning for Leduc Hold'em poker — a tractable benchmark with a computable Nash equilibrium. Standard RL agents learn a fixed policy via self-play and do not adapt to opponent tendencies. We propose augmenting a Q-learning agent with a lightweight opponent model that tracks action frequency statistics (bet, raise, fold rates) over recent hands and classifies the opponent into one of three behavioral styles: random, tight, or aggressive. The Q-learning agent's state representation is extended with this style indicator, enabling it to adapt its strategy accordingly. We use the RLCARD toolkit [12] for the Leduc Hold'em environment and baseline agents, ensuring reproducible experiments. Baselines include a random agent, a rule-based agent, and a standard Q-learning agent. We evaluate using win rate, average reward, and learning curves.

Index Terms—reinforcement learning, poker, Leduc Hold'em, opponent modeling, Q-learning, imperfect information games

I. INTRODUCTION

Poker is a canonical benchmark for decision-making under uncertainty. Unlike perfect-information games such as Chess or Go, it requires reasoning over hidden state, managing multi-round strategy, and adapting to opponent behavior — all core challenges in reinforcement learning (RL).

Recent breakthroughs such as Libratus [1] and Pluribus [2] achieved superhuman performance in No-Limit Texas Hold'em (NLHE), but these systems primarily target Nash equilibrium strategies that are robust against any opponent rather than adaptive strategies that exploit specific tendencies. Moreover, NLHE's enormous game tree ($\sim 10^{164}$ states) makes it impractical for course-level research. Leduc Hold'em [8], with only 6 cards and 2 betting rounds, preserves the core challenges of poker — hidden information, bluffing, and multi-round reasoning — while remaining computationally tractable.

Most RL agents for poker assume a stationary opponent. In practice, opponents exhibit consistent behavioral patterns: some fold frequently (tight), others bet and raise aggressively, and others play randomly. An agent that detects and exploits these patterns should outperform a fixed-policy baseline. However, this hypothesis has not been studied in a controlled, reproducible setting on Leduc Hold'em. Our contributions are:

- A simple, frequency-based opponent style classifier requiring no recurrent networks or Bayesian inference.
- An opponent-aware Q-learning agent that conditions its policy on the inferred opponent style.

- A comparison against random, rule-based, and standard Q-learning baselines using win rate and average reward.

II. LITERATURE SURVEY

A. Counterfactual Regret Minimization (CFR)

CFR [5] is the foundational algorithm for solving extensive-form imperfect-information games. By iteratively minimizing counterfactual regret at each information set, it converges to a Nash equilibrium. Variants including CFR+ [6] and Monte Carlo CFR [7] improve practical convergence speed. RLCARD provides a built-in CFR implementation for Leduc Hold'em, which we use directly as an approximate Nash reference agent.

B. Neural Fictitious Self-Play (NFSP)

NFSP [3] combines DQN-based RL with a supervised average-strategy network, provably converging to Nash equilibrium in two-player zero-sum games. While powerful, NFSP requires maintaining two separate networks and a reservoir buffer, making it complex to implement from scratch. We use RLCARD's NFSP implementation as a reference point in our literature survey but simplify our own baseline to Q-learning.

C. Superhuman Agents in NLHE

DeepStack [4] introduced depth-limited solving with a learned value network, achieving human-expert level in heads-up NLHE. Libratus [1] combined abstraction, real-time solving, and opponent exploitation to defeat professional players. Pluribus [2] extended this to six-player NLHE. These systems target Nash equilibrium rather than explicit opponent adaptation, and their scale makes them unsuitable as direct baselines for our work.

D. Opponent Modeling in Poker

Early opponent modeling used Bayesian methods to maintain distributions over opponent hand ranges [8]. More recent work encodes opponent observations into neural networks using Mixture-of-Experts architectures to discover opponent strategy clusters [10]. Rahman et al. [11] proposed AMP3 for multi-player Texas Hold'em, feeding predicted opponent style features into an Actor-Critic agent. However, all of these approaches target NLHE, where win rate is the only available metric. Leduc Hold'em allows a more rigorous comparison since the Nash equilibrium is known. To our knowledge, no prior work has studied frequency-based opponent modeling specifically on Leduc Hold'em with Q-learning as the base agent.

III. BACKGROUND

A. Leduc Hold'em

Leduc Hold'em is a two-player poker game with a 6-card deck (Jack, Queen, King in two suits). Each player receives one private card and posts a fixed ante of 1 chip. Play proceeds over two betting rounds: pre-flop and post-flop (after one community card is revealed). Each round permits a maximum of two raises. The winner is determined by showdown rank (pairs beat high cards) or by the opponent folding.

B. Problem Formulation

We model Leduc Hold'em as a Partially Observable Markov Decision Process (POMDP), formally defined by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{O})$:

TABLE I
POMDP COMPONENTS FOR LEDUC HOLD'EM

Component	Definition
State $s \in \mathcal{S}$	Full game state: both players' cards, community card, pot size, betting history
Observation $o \in \mathcal{O}$	Agent's private view: $o = (\text{private_card}, \text{public_card}, \text{pot_size}, \text{betting_history})$
Actions $\mathcal{A}$	$\mathcal{A} = \{\text{fold}, \text{call}, \text{raise}\}$
Transition $\mathcal{T}$	Stochastic: determined by opponent's hidden card and betting responses
Reward $r \in \mathcal{R}$	Net chip change at end of hand; $r = 0$ at all intermediate steps

Since the opponent's private card is never revealed (unless showdown occurs), the agent operates on observations o rather than full states s . In practice, RLCard encodes o as a fixed-length vector, and the state space is small enough to support tabular Q-learning without function approximation.

IV. BASELINE IMPLEMENTATION

We use the RLCard toolkit [12] for all experiments. RLCard provides the Leduc Hold'em environment, a pretrained CFR agent as an approximate Nash baseline, and evaluation utilities — eliminating the need to implement game logic from scratch.

Random Agent. Selects uniformly among legal actions. Serves as the lower-bound baseline.

Rule-Based Agent. A hand-coded heuristic agent: raise with a pair, call with a high card (King), fold otherwise. Represents a simple but non-trivial fixed strategy.

Q-Learning Agent. A tabular Q-learning agent trained against a random opponent using ϵ -greedy exploration. This is our primary RL baseline and the foundation for the proposed method.

V. PROPOSED METHOD: OPPONENT-AWARE Q-LEARNING

A. Motivation

The standard Q-learning agent treats the opponent as a fixed, unknown part of the environment. However, real opponents exhibit systematic behavioral patterns. An agent that detects these patterns and adapts its strategy should achieve higher

win rates than one that ignores them. We propose a minimal, practical mechanism to achieve this without recurrent networks or Bayesian inference.

B. Opponent Style Classification

We track three action frequency statistics over a sliding window of the last $W = 20$ observed opponent actions:

$$f_{\text{bet}} = \frac{N_{\text{bet}}}{W}, \quad f_{\text{raise}} = \frac{N_{\text{raise}}}{W}, \quad f_{\text{fold}} = \frac{N_{\text{fold}}}{W} \quad (1)$$

The opponent is classified into one of three styles using simple thresholds:

- **Aggressive:** $f_{\text{bet}} + f_{\text{raise}} > 0.6$
- **Tight:** $f_{\text{fold}} > 0.6$
- **Random:** otherwise (default at game start)

This requires no learned parameters and updates in $O(1)$ per step.

C. Augmented State Representation

The opponent-aware agent's state is augmented with the inferred style indicator $c \in \{\text{random}, \text{tight}, \text{aggressive}\}$:

$$s' = (s, c) \quad (2)$$

where s is the standard RLCard observation vector. The Q-table is indexed by (s', a) , allowing the agent to learn distinct policies for each opponent style. No neural network or sequence model is required.

D. Training Against Fixed-Style Opponents

We train the opponent-aware agent against three scripted opponents:

- **Random opponent:** selects actions uniformly.
- **Tight opponent:** folds unless holding a pair or King; never raises.
- **Aggressive opponent:** raises whenever possible; folds only with Jack pre-flop.

Each training episode samples one opponent type uniformly. The agent learns to classify the opponent from its observed actions and adapt its strategy accordingly.

E. Evaluation

We evaluate on win rate and average reward per episode across 10,000 evaluation games, comparing:

TABLE II
AGENT COMPARISON SUMMARY

Agent	Purpose
Random	Lower bound
Rule-based	Heuristic baseline
Q-learning	RL baseline (style-unaware)
Opponent-aware Q-learning	Proposed method
CFR (RLCard)	Approximate Nash reference

We report win rate and average reward against each fixed-style opponent separately and in aggregate. Learning curves (reward vs. training episodes) are reported for both Q-learning and the proposed method.

VI. DISCUSSION AND CONCLUSION

Our central hypothesis is that the opponent-aware agent achieves higher win rates against stylistically biased opponents compared to the style-unaware Q-learning baseline, particularly against tight and aggressive opponents where the behavioral signal is strong. Against a random opponent, both agents should perform similarly since no exploitable pattern exists.

A potential risk is that style classification is noisy at the start of an episode (before enough actions are observed), which may lead to suboptimal early decisions. We will analyze performance as a function of episode progress to characterize this warmup effect.

In summary, we propose a lightweight, practical approach to opponent-aware RL in Leduc Hold'em. By using action frequency statistics rather than recurrent networks, and Q-learning rather than NFSP, our method is implementable within a course project scope (estimated ~ 600 lines of Python) while still addressing a genuine open question: does simple opponent style detection meaningfully improve RL agent performance in a tractable poker environment? Full experimental results will be presented in the final report.

REFERENCES

- [1] N. Brown and T. Sandholm, "Libratus: The Superhuman AI for No-Limit Poker," in *Proc. IJCAI*, 2017, pp. 5226–5228.
- [2] N. Brown and T. Sandholm, "Superhuman AI for multiplayer poker," *Science*, vol. 365, no. 6456, pp. 885–890, 2019.
- [3] J. Heinrich and D. Silver, "Deep Reinforcement Learning from Self-Play in Imperfect-Information Games," *arXiv preprint arXiv:1603.01121*, 2016.
- [4] M. Moravčík *et al.*, "DeepStack: Expert-level artificial intelligence in heads-up no-limit poker," *Science*, vol. 356, no. 6337, pp. 508–513, 2017.
- [5] M. Zinkevich, M. Johanson, M. Bowling, and C. Piccione, "Regret minimization in games with incomplete information," in *Proc. NeurIPS*, 2007, pp. 1729–1736.
- [6] O. Tammelin, "Solving large imperfect information games using CFR+," *arXiv preprint arXiv:1407.5042*, 2014.
- [7] M. Lanctot, K. Waugh, M. Zinkevich, and M. Bowling, "Monte Carlo sampling for regret minimization in extensive games," in *Proc. NeurIPS*, 2009.
- [8] F. Southey *et al.*, "Bayes' Bluff: Opponent Modelling in Poker," in *Proc. UAI*, 2005.
- [9] S. Ganzfried and T. Sandholm, "Endgame Solving in Large Imperfect-Information Games," in *Proc. AAMAS*, 2015.
- [10] C. Hsieh *et al.*, "Opponent Modeling with Neural Architecture for Poker," *arXiv preprint*, 2023.
- [11] A. Rahman *et al.*, "AMP3: Adaptive Multi-Player Poker via Opponent Style Modeling," *arXiv preprint*, 2024.
- [12] D. Zha *et al.*, "RLCard: A Toolkit for Reinforcement Learning in Card Games," in *Proc. IJCAI*, 2020.